

Implementing Feed-forward neural networks with Keras and TensorFlow

-Tushar Bhagat Roll_No. 07

```
In [23]: import numpy as np
import matplotlib.pyplot as plt
import tensorflow as tf
from tensorflow.keras.datasets import mnist
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense,Input
from tensorflow.keras.optimizers import SGD
from tensorflow.keras.utils import to_categorical
from tensorflow.keras.utils import plot_model

In [24]: # Load MNIST data
(x_train, y_train), (x_test, y_test) = mnist.load_data()

# Preprocess data
x_train = x_train.reshape((x_train.shape[0], 28 * 28)).astype('float32') / 255
x_test = x_test.reshape((x_test.shape[0], 28 * 28)).astype('float32') / 255

# One-hot encode labels
y_train = to_categorical(y_train, 10)
y_test = to_categorical(y_test, 10)

In [25]: # Define the network architecture
model = Sequential([
    Input(shape=(28 * 28,)), # Define the input shape here
    Dense(512, activation='relu'),
    Dense(256, activation='relu'),
    Dense(10, activation='softmax')
])

In [26]: # Compile the model
model.compile(optimizer=SGD(),
              loss='categorical_crossentropy',
              metrics=['accuracy'])

# Train the model
history = model.fit(x_train, y_train,
                   epochs=20,
                   batch_size=32,
                   validation_split=0.2)
```

Epoch 1/20
1500/1500 ————— 5s 3ms/step - accuracy: 0.7532 - loss: 0.9979 - val_accuracy: 0.9118 - val_loss: 0.3131
Epoch 2/20
1500/1500 ————— 5s 3ms/step - accuracy: 0.9129 - loss: 0.3045 - val_accuracy: 0.9302 - val_loss: 0.2435
Epoch 3/20
1500/1500 ————— 5s 4ms/step - accuracy: 0.9306 - loss: 0.2419 - val_accuracy: 0.9386 - val_loss: 0.2181
Epoch 4/20
1500/1500 ————— 5s 3ms/step - accuracy: 0.9401 - loss: 0.2056 - val_accuracy: 0.9479 - val_loss: 0.1887
Epoch 5/20
1500/1500 ————— 5s 3ms/step - accuracy: 0.9477 - loss: 0.1812 - val_accuracy: 0.9531 - val_loss: 0.1687
Epoch 6/20
1500/1500 ————— 4s 3ms/step - accuracy: 0.9556 - loss: 0.1618 - val_accuracy: 0.9563 - val_loss: 0.1547
Epoch 7/20
1500/1500 ————— 4s 3ms/step - accuracy: 0.9600 - loss: 0.1420 - val_accuracy: 0.9594 - val_loss: 0.1429
Epoch 8/20
1500/1500 ————— 5s 3ms/step - accuracy: 0.9635 - loss: 0.1303 - val_accuracy: 0.9628 - val_loss: 0.1337
Epoch 9/20
1500/1500 ————— 6s 4ms/step - accuracy: 0.9684 - loss: 0.1133 - val_accuracy: 0.9647 - val_loss: 0.1267
Epoch 10/20
1500/1500 ————— 5s 3ms/step - accuracy: 0.9719 - loss: 0.1036 - val_accuracy: 0.9668 - val_loss: 0.1183
Epoch 11/20
1500/1500 ————— 5s 4ms/step - accuracy: 0.9740 - loss: 0.0969 - val_accuracy: 0.9684 - val_loss: 0.1131
Epoch 12/20
1500/1500 ————— 5s 3ms/step - accuracy: 0.9753 - loss: 0.0889 - val_accuracy: 0.9680 - val_loss: 0.1118
Epoch 13/20
1500/1500 ————— 5s 3ms/step - accuracy: 0.9788 - loss: 0.0813 - val_accuracy: 0.9693 - val_loss: 0.1055
Epoch 14/20
1500/1500 ————— 5s 3ms/step - accuracy: 0.9802 - loss: 0.0765 - val_accuracy: 0.9705 - val_loss: 0.1027
Epoch 15/20
1500/1500 ————— 6s 4ms/step - accuracy: 0.9822 - loss: 0.0690 - val_accuracy: 0.9721 - val_loss: 0.0981
Epoch 16/20
1500/1500 ————— 6s 4ms/step - accuracy: 0.9829 - loss: 0.0634 - val_accuracy: 0.9718 - val_loss: 0.0967
Epoch 17/20
1500/1500 ————— 6s 4ms/step - accuracy: 0.9843 - loss: 0.0608 - val_accuracy: 0.9722 - val_loss: 0.0958
Epoch 18/20
1500/1500 ————— 5s 3ms/step - accuracy: 0.9854 - loss: 0.0569 - val_accuracy: 0.9738 - val_loss: 0.0903
Epoch 19/20
1500/1500 ————— 5s 3ms/step - accuracy: 0.9858 - loss: 0.0538 - val_accuracy: 0.9744 - val_loss: 0.0892
Epoch 20/20
1500/1500 ————— 6s 4ms/step - accuracy: 0.9872 - loss: 0.0489 - val_accuracy: 0.9733 - val_loss: 0.0933

```
In [27]: model.summary()
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
dense_6 (Dense)	(None, 512)	401,920
dense_7 (Dense)	(None, 256)	131,328
dense_8 (Dense)	(None, 10)	2,570

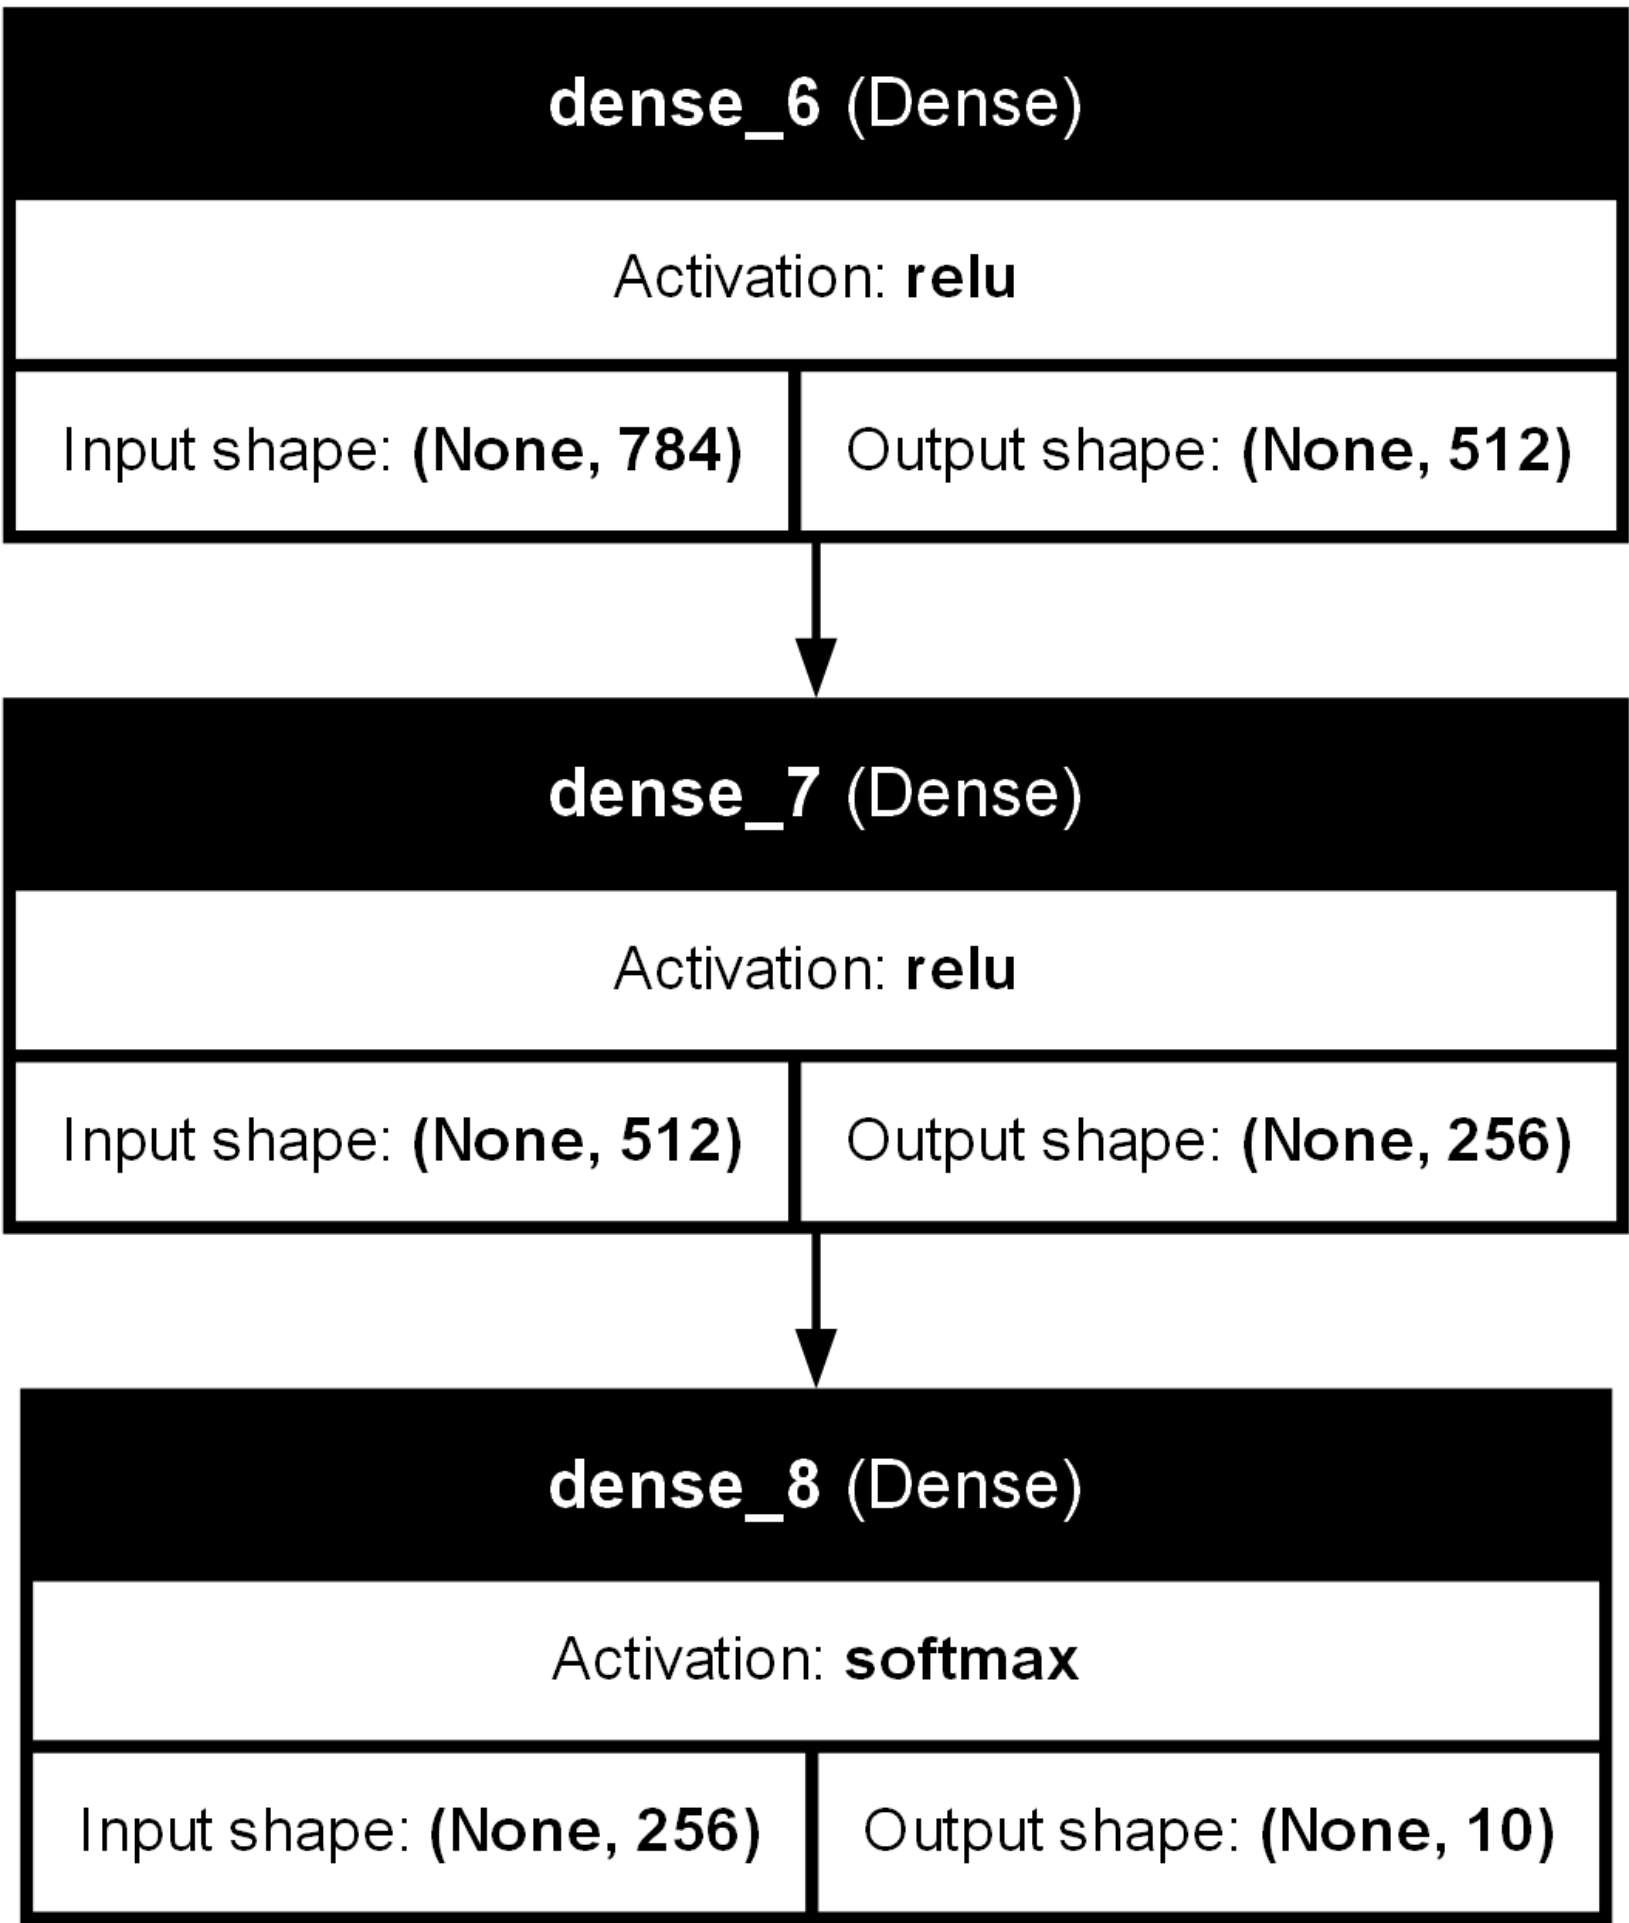
Total params: 535,820 (2.04 MB)
Trainable params: 535,818 (2.04 MB)
Non-trainable params: 0 (0.00 B)
Optimizer params: 2 (12.00 B)

```
In [28]: # Evaluate the model on test data
test_loss, test_accuracy = model.evaluate(x_test, y_test)
print(f'Test Loss: {test_loss:.4f}')
print(f'Test Accuracy: {test_accuracy:.4f}')
```

313/313 ————— 1s 2ms/step - accuracy: 0.9695 - loss: 0.1002
Test Loss: 0.0863
Test Accuracy: 0.9732

```
In [29]: #plotting the model
plot_model(model,to_file='C:/Users/tmbha/Downloads/DLprac2_model.png',show_shapes=True,show_layer_names=True,show_l
```

Out[29]:



```
In [30]: # Plotting
plt.figure(figsize=(14, 6))

# Plot training & validation Loss
plt.subplot(1, 2, 1)
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title('Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()

# Plot training & validation accuracy
plt.subplot(1, 2, 2)
plt.plot(history.history['accuracy'], label='Train Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.title('Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()

plt.show()
```

